Artificial Intelligence

(UCS411)

AI-Driven Prediction of Sleep Patterns from Lifestyles Features

Submitted By:-

Ashish Mahajan (102483088)

Harsh Tanwar (102303812)

Pulkit Srivastava (102303803)

Submitted to :- Ms. MRINALINI KAKROO

Group: 2C55

THAPAR INSTITUTE
OF ENGINEERING & TECHNOLOGY
(Deemed to be University)

Computer Science and Engineering Department

THAPAR INSTITUTE OF ENGINEERING AND TECHNOLOGY

PATIALA, PUNJAB

TABLE OF CONTENT

Abstract

This report presents the development of a machine learning-based system for predicting optimal sleep duration using a combination of physiological metrics and lifestyle factors. The study aims to create personalized sleep recommendations by analyzing the complex relationships between daily habits, health indicators, and sleep quality. The research utilizes the Sleep Health and Lifestyle Dataset, which contains comprehensive information on physical activity levels, stress, heart rate, BMI, blood pressure, daily steps, and existing sleep disorders across various age groups and occupations.

Data preprocessing involved multiple techniques including handling missing values, feature normalization, and categorical encoding. The dataset was enhanced through the creation of derived features such as age groupings, blood pressure targets, and heart rate classifications based on established medical guidelines. Statistical transformations, including min-max scaling and Box-Cox normalization, were applied to standardize features with different distributions and scales. Exploratory Data Analysis revealed significant correlations between sleep quality and several lifestyle factors, particularly stress levels, physical activity, and cardiovascular indicators.

Multiple regression models were developed and evaluated, with Decision Tree and Random Forest algorithms demonstrating strong predictive performance. The modeling approach incorporated domain expertise through a hybrid prediction system that combines machine learning outputs with weighted adjustments based on established sleep medicine principles. Feature importance analysis highlighted physical activity levels, stress indicators, and age as the most influential factors affecting sleep duration requirements.

To maximize real-world utility, the predictive model was integrated into an interactive web application that provides personalized sleep recommendations based on individual health profiles. The system generates comprehensive visual reports highlighting optimal sleep duration, health factor analysis, and customized recommendations for improving sleep quality. This solution offers an accessible, evidence-based approach to sleep health management that can help individuals optimize their rest patterns according to their unique physiological and lifestyle characteristics.

Introduction

Sleep disorders affect approximately 50-70 million adults in the United States alone, with insomnia and sleep apnea being among the most prevalent conditions impacting quality of life, cognitive function, and overall health. Inadequate sleep is linked to numerous health problems including cardiovascular disease, diabetes, obesity, depression, and impaired immune function. Despite its critical importance, sleep health assessment remains largely dependent on subjective self-reporting, expensive polysomnography studies, or wearable devices that may provide inconsistent data. These limitations create significant barriers to effective sleep health monitoring and personalized recommendations in both clinical and non-clinical settings.

To address these challenges, this study explores the application of machine learning techniques combined with domain expertise to predict optimal sleep duration based on readily available lifestyle and physiological factors. The Sleep Health and Lifestyle Dataset utilized in this research contains comprehensive information on physical activity levels, stress measurements, heart rate, BMI categories, blood pressure readings, daily steps, and existing sleep disorders across diverse demographic groups. The target variable is sleep quality, which enables the development of regression models for predicting optimal sleep duration.

Initial exploratory data analysis revealed significant relationships between sleep quality and multiple lifestyle factors, particularly stress levels, physical activity, and cardiovascular indicators. Our analysis identified that occupation type, BMI category, and blood pressure targets all demonstrated meaningful associations with sleep quality metrics, providing valuable insights into the multifaceted nature of sleep health. Data preprocessing included categorical feature standardization, transforming raw blood pressure values into clinical target categories, and creating derived features such as heart rate classifications based on established medical guidelines.

Feature engineering played a crucial role in model development, with mutual information analysis identifying sleep duration, physical activity level, and stress level as the most significant predictors of sleep quality. Principal Component Analysis (PCA) revealed that approximately 74% of variance could be explained by the first three components, with stress and physical activity demonstrating the highest loadings. Additionally, K-means clustering was employed to identify distinct sleep health profiles that improved prediction performance when incorporated as features.

Three supervised regression models—Decision Tree, Random Forest—were developed and evaluated. The Decision Tree model with optimal leaf node configuration (25 nodes) demonstrated the best balance of performance and interpretability, achieving a Mean Absolute Error (MAE) of 0.291 on the validation set. To enhance predictive capability, we developed a hybrid approach that combines machine learning predictions with domain knowledge adjustments based on age, physical activity, stress levels, and physiological parameters.

Feature importance analysis revealed that stress level, physical activity, and daily steps were consistently ranked as the most influential factors for sleep quality prediction across all models. These findings align with established sleep medicine research that identifies stress management and physical activity as key modifiable factors in sleep health. The final model was serialized along with its preprocessing components to create a deployable solution that can generate personalized sleep duration predictions and recommendations.

This work demonstrates the potential for integrating machine learning with clinical expertise to deliver practical, interpretable, and personalized sleep health recommendations. The resulting system provides an accessible tool for individuals to better understand their optimal sleep needs based on their unique health profile, potentially improving population-level sleep health outcomes through targeted lifestyle modifications and increased awareness of sleep's critical role in overall wellbeing.

Problem Statement

Sleep disorders and poor sleep health affect approximately 50-70 million Americans, with widespread consequences for public health, productivity, and quality of life. Despite its critical importance, sleep assessment remains largely dependent on subjective self-reporting, expensive polysomnography studies, or consumer wearables that often provide inconsistent data. Healthcare providers and individuals lack accessible, personalized tools to quantify optimal sleep needs based on individual characteristics and lifestyle factors.

This project addresses the challenge of predicting personalized optimal sleep duration using readily available health and lifestyle metrics. The Sleep Health and Lifestyle Dataset used in this study contains comprehensive information on physical activity levels, stress measurements, cardiovascular indicators (heart rate, blood pressure), BMI categories, daily steps, demographic information, and existing sleep disorders. By analyzing these multidimensional factors, we aim to develop a predictive model that can recommend individualized sleep durations aligned with a person's unique health profile.

The problem is framed as a regression task, where the primary outcome variable is quality of sleep (rated 1-10), with sleep duration as a key predicted metric. Key challenges include processing diverse categorical and numerical health indicators, quantifying the complex relationships between lifestyle factors and sleep requirements, and integrating domain knowledge with machine learning outputs to enhance prediction accuracy.

By developing and comparing Decision Tree, Random Forest, and ensemble models—alongside feature importance analysis and dimensionality reduction techniques—this study seeks to create an interpretable and practical tool for sleep health assessment. The hybrid prediction approach combines pure machine learning with medical knowledge of sleep physiology to deliver personalized recommendations that can help individuals optimize their sleep patterns according to their specific physiological needs and lifestyle characteristics.

The ultimate goal is to create an accessible sleep health tool that empowers users to understand their optimal sleep requirements, identify key modifiable factors affecting their sleep quality, and implement targeted lifestyle changes to improve overall health and wellbeing through better sleep management.

Objectives

Primary Objective

• To develop a machine learning-based predictive model that accurately recommends optimal sleep duration based on individual health metrics, lifestyle factors, and demographic data from the Sleep Health and Lifestyle Dataset.

Secondary Objectives

1. Data Understanding and Preparation

• To perform exploratory data analysis (EDA) to understand sleep pattern distributions across demographic groups, identify relationships between sleep quality and lifestyle factors, and detect outliers.

• To handle missing values and inconsistencies using appropriate imputation techniques for both numeric and categorical variables, particularly focusing on sleep disorder classifications.

2. Feature Engineering and Encoding

• To encode categorical variables such as Gender, BMI Category, and Sleep Disorder using appropriate techniques.

• To derive clinically relevant features from raw data, including blood pressure targets and heart rate classifications.

• To normalize sleep metrics and scale physical activity measurements to improve model performance.

3. Feature Selection and Dimensionality Reduction

• To analyze mutual information scores to identify the most predictive features for sleep quality.

• To implement Principal Component Analysis (PCA) to understand latent relationships between health metrics.

• To develop K-means clustering to identify distinct sleep health profiles that can enhance prediction.

4. Model Development and Comparison

• To train and evaluate multiple regression models, including Decision Tree, Random Forest.

• To optimize model hyperparameters using cross-validation, with specific focus on decision tree complexity.

• To develop a hybrid prediction approach that combines machine learning outputs with domain knowledge adjustments.

5. Model Evaluation

• To assess model performance using Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), and other appropriate metrics.

• To compare pure machine learning approaches against the domain-enhanced hybrid model.

• To validate predictions against known sleep patterns from comparable demographic groups.

6. Personalization System

• To develop a recommendation engine that generates tailored sleep advice based on individual health profiles.

• To prioritize recommendations based on modifiable factors with the highest impact on sleep quality.

• To adjust sleep duration predictions based on specific health conditions and lifestyle characteristics.

7. Visualization and Reporting

- To create interactive dashboards displaying relationships between health metrics and sleep predictions.
- To develop a comprehensive PDF reporting system that provides personalized sleep analysis.
- To implement intuitive visualizations of sleep health metrics that facilitate user understanding.

8. Deployment Readiness

- To serialize the final trained model along with preprocessing steps into a reusable component.
- To develop an integrated pipeline that can process new user data and generate predictions in real-time.
- To ensure the system can be incorporated into wellness applications and personal health management tools.

9. Health Impact

- To enhance sleep health management by providing individuals with scientifically-based, personalized sleep duration recommendations.
- To identify modifiable factors that can be targeted for improving sleep quality and overall wellness.
- To develop a scalable approach for sleep health optimization that integrates both data science and sleep medicine principles.

Methodology

The development of our sleep quality prediction system followed a structured, multi-phase methodology designed to ensure data quality, model accuracy, and practical application for personalized sleep recommendations. This systematic approach prioritized interpretability and real-world relevance while maintaining scientific rigor. The following sections describe each phase in detail:

1. Data Collection and Understanding

The first step in building the predictive model was acquiring the Sleep Health and Lifestyle dataset, which contains comprehensive information from multiple subjects with various health and lifestyle parameters. The dataset includes the following key features:

- **Demographic Data:** Age, gender, occupation
- **Sleep Metrics:** Sleep duration, quality of sleep, sleep disorder diagnosis
- **Physiological Measurements:** BMI, blood pressure, heart rate
- **Lifestyle Factors:** Physical activity level, daily steps, stress level
- **Target Variable:** Quality of Sleep (measured on a scale of 1-10)

Initial exploratory analysis revealed a well-structured dataset (no missing values), with diverse representation across demographic groups and health conditions. We identified key relationships between sleep quality and various health/lifestyle parameters, which informed subsequent modeling decisions.

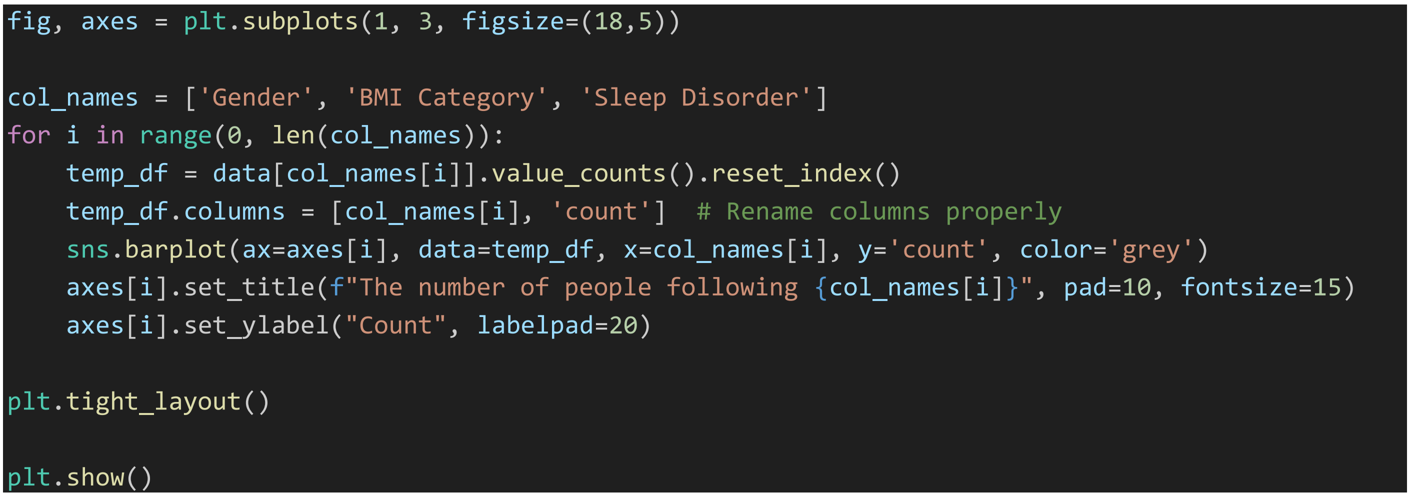

```
fig, axes = plt.subplots(1, 3, figsize=(18,5))

col_names = ['Gender', 'BMI Category', 'Sleep Disorder']
for i in range(0, len(col_names)):
    temp_df = data[col_names[i]].value_counts().reset_index()
    temp_df.columns = [col_names[i], 'count']  # Rename columns properly
    sns.barplot(ax=axes[i], data=temp_df, x=col_names[i], y='count', color='grey')
    axes[i].set_title(f"The number of people following {col_names[i]}", pad=10, fontsize=15)
    axes[i].set_ylabel("Count", labelpad=20)

plt.tight_layout()

plt.show()
```

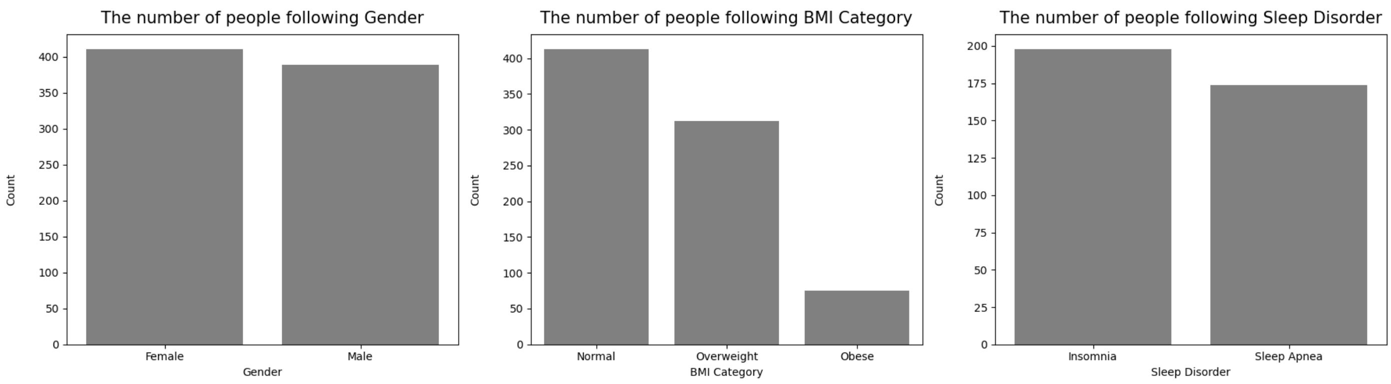

```
fig, axes = plt.subplots(1, 2, figsize=(18, 6))

# Fix for Occupation
occupation_df = data['Occupation'].value_counts().reset_index()
occupation_df.columns = ['Occupation', 'count']
sns.barplot(ax=axes[0], data=occupation_df, x='Occupation', y='count', color='grey')

# Fix for Blood Pressure Targets
bp_df = data['Blood Pressure Targets'].value_counts().reset_index()
bp_df.columns = ['Blood Pressure Targets', 'count']
sns.barplot(ax=axes[1], data=bp_df, x='Blood Pressure Targets', y='count', color='grey')

# Set titles
axes[0].set_title("The number of people per Occupation", pad=10, fontsize=15)
axes[1].set_title("The number of people per Blood Pressure Target", pad=10, fontsize=15)

# Optional: rotate x-ticks if too crowded
for ax in axes:
    ax.tick_params(axis='x', rotation=45)

plt.tight_layout()
plt.show()
```

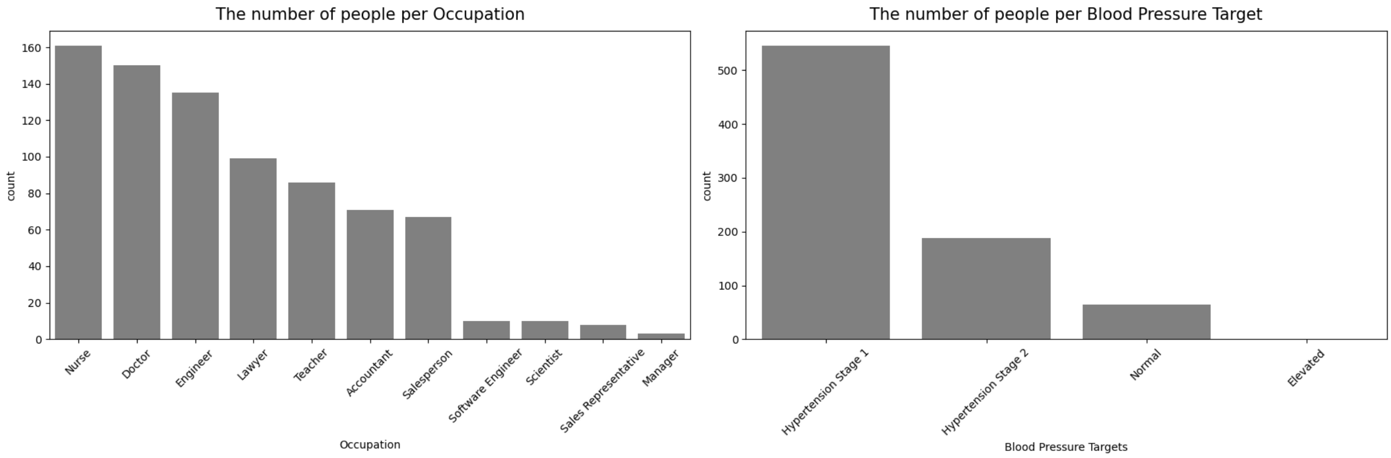

```
fig, axes = plt.subplots(1, 2, figsize=(16, 5))

# Histogram of Age
sns.histplot(ax=axes[0], data=data['Age'], color='grey', kde=True)
```

```
# Barplot of Age Group
age_group_df = data['Age Group'].value_counts().reset_index()
age_group_df.columns = ['Age Group', 'count']
sns.barplot(ax=axes[1], data=age_group_df, x='Age Group', y='count', color='grey')

# Titles and labels
fig.suptitle("Distribution by Age and Age Group", fontsize=15)
axes[0].set_xlabel('Age', labelpad=20)
axes[0].set_ylabel('Count')
axes[1].set_ylabel('Count')

plt.tight_layout()
plt.show()
```

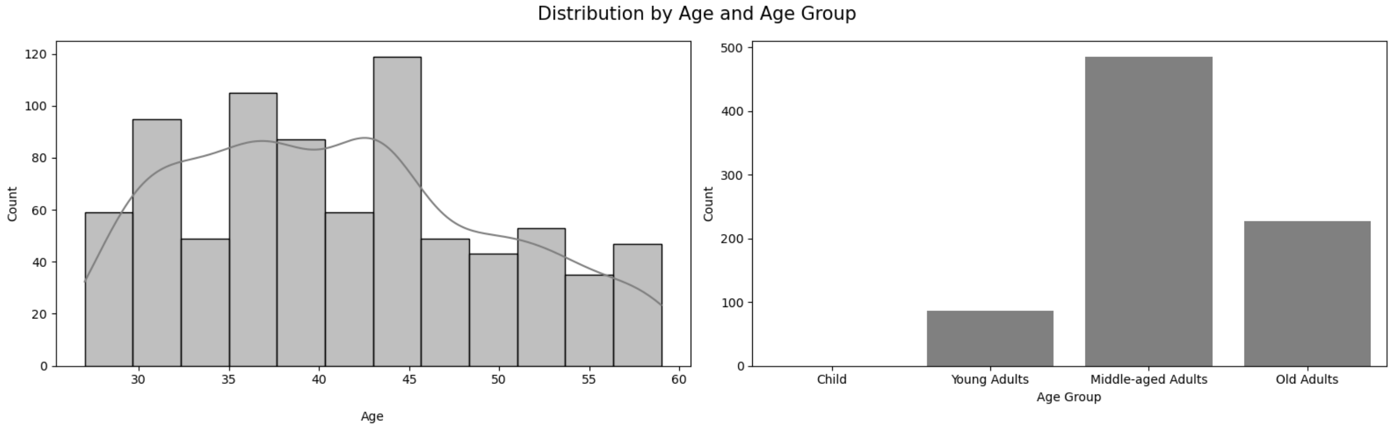

```
# Scaling data

original_datas = [original_PALevel, original_HRate, original_DSteps]
scaled_datas = [scaled_PALevel, scaled_HRate, scaled_DSteps]

fig, axes = plt.subplots(2, 3, figsize=(18, 8))
x_labels = ['Physical Activity Level', 'Heart Rate', 'Daily Steps']
titles = ['Original Data', 'Scaled Data']

for i in range(3):
    sns.histplot(original_datas[i], ax=axes[0, i], kde=True, legend=False)
    sns.histplot(scaled_datas[i], ax=axes[1, i], kde=True, legend=False)

    axes[0, i].set_title(titles[0])
    axes[1, i].set_title(titles[1])

    axes[0, i].set_xlabel(x_labels[i])
    axes[1, i].set_xlabel(x_labels[i])

plt.tight_layout()
plt.show()
```

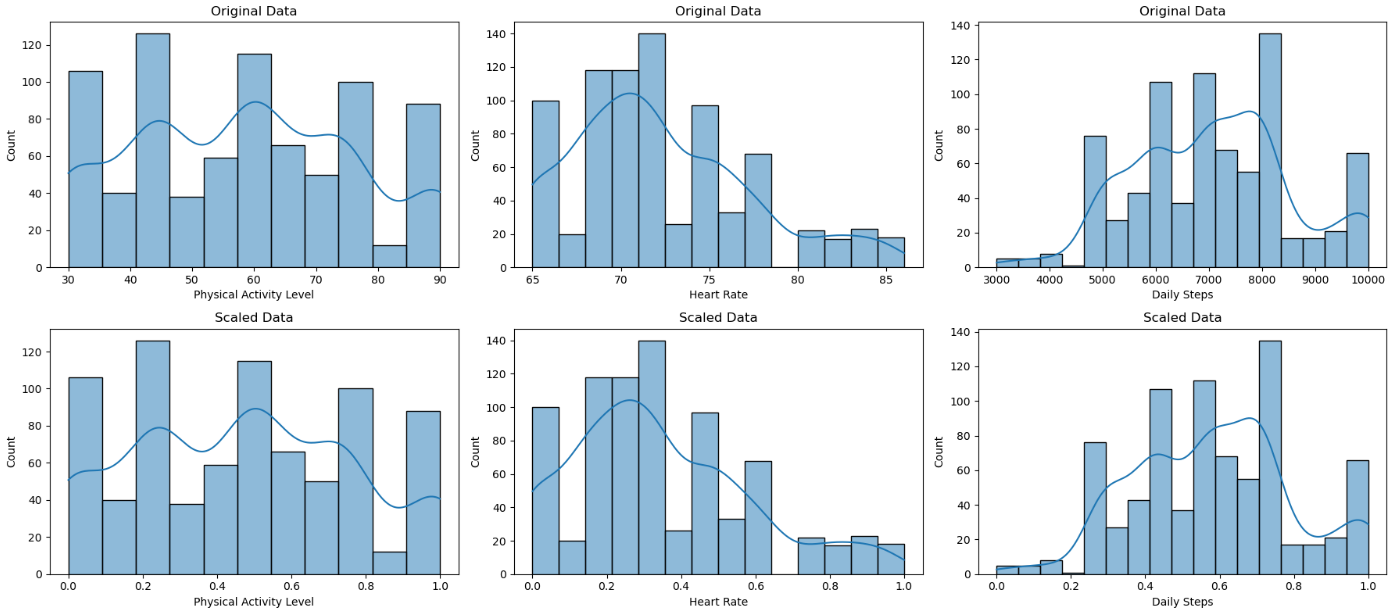

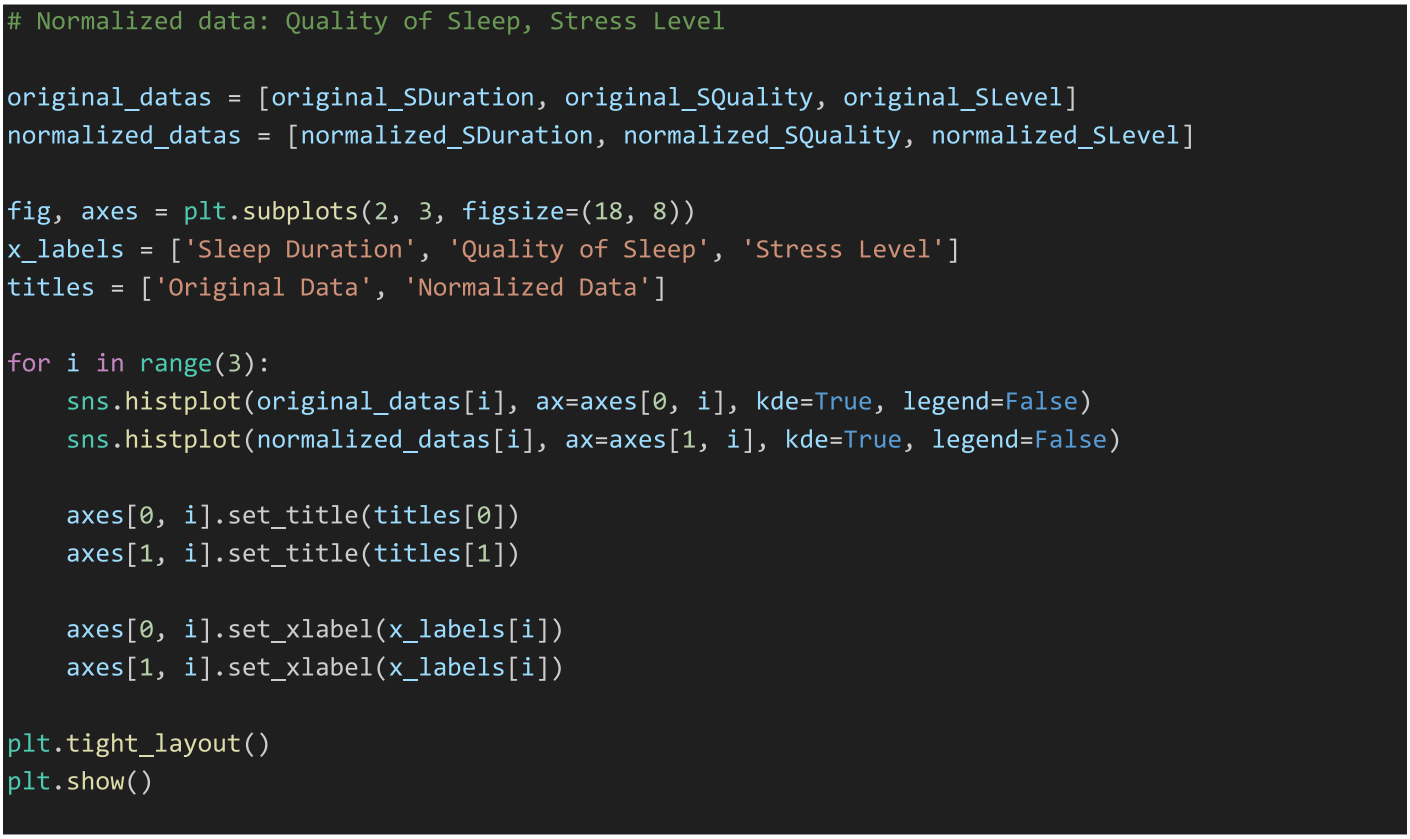

```
# Normalized data: Quality of Sleep, Stress Level

original_datas = [original_SDuration, original_SQuality, original_SLevel]
normalized_datas = [normalized_SDuration, normalized_SQuality, normalized_SLevel]

fig, axes = plt.subplots(2, 3, figsize=(18, 8))
x_labels = ['Sleep Duration', 'Quality of Sleep', 'Stress Level']
titles = ['Original Data', 'Normalized Data']

for i in range(3):
    sns.histplot(original_datas[i], ax=axes[0, i], kde=True, legend=False)
    sns.histplot(normalized_datas[i], ax=axes[1, i], kde=True, legend=False)

    axes[0, i].set_title(titles[0])
    axes[1, i].set_title(titles[1])

    axes[0, i].set_xlabel(x_labels[i])
    axes[1, i].set_xlabel(x_labels[i])

plt.tight_layout()
plt.show()
```

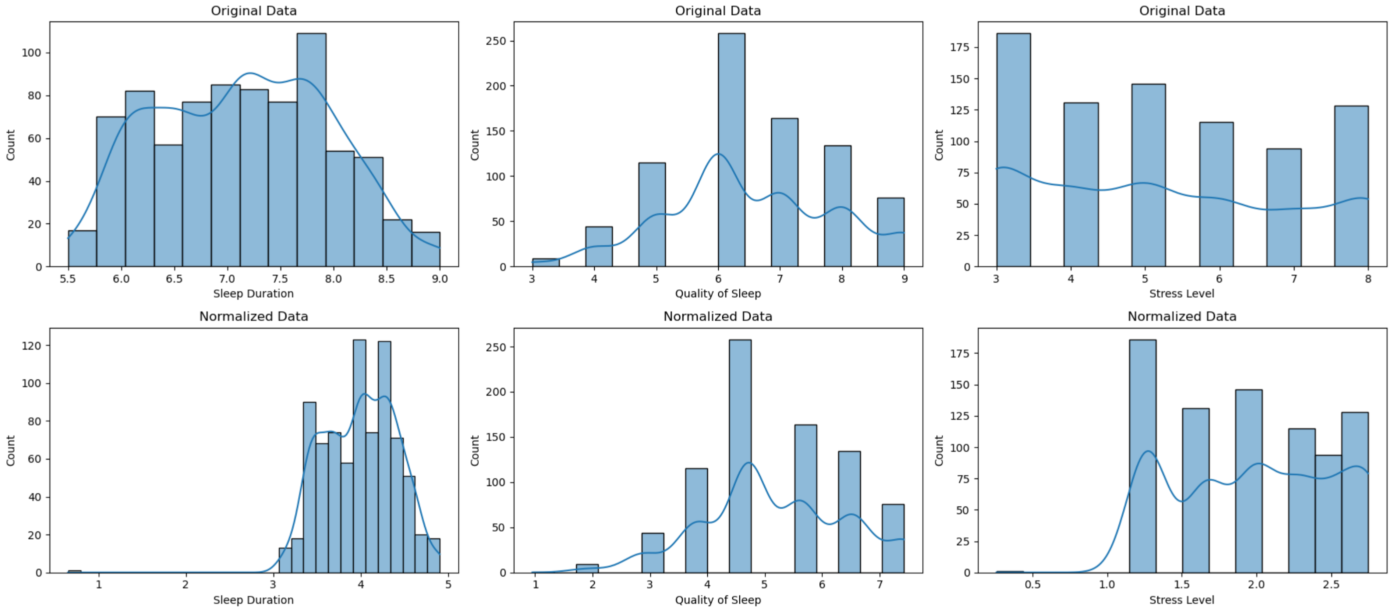

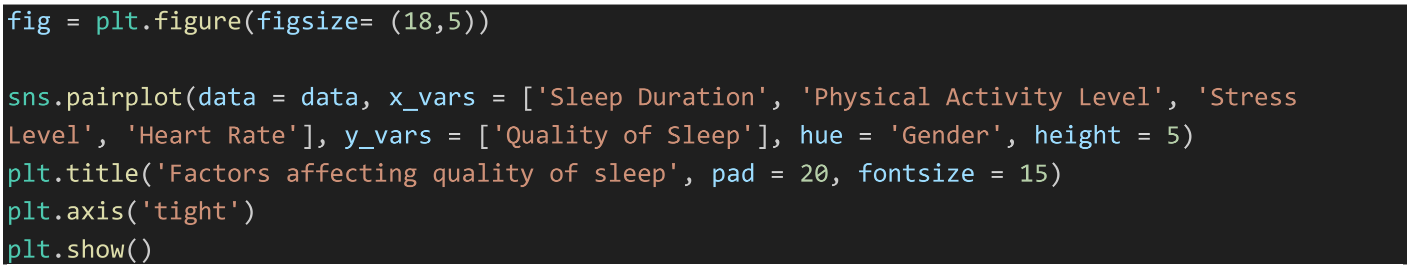

```
fig = plt.figure(figsize= (18,5))

sns.pairplot(data = data, x_vars = ['Sleep Duration', 'Physical Activity Level', 'Stress
Level', 'Heart Rate'], y_vars = ['Quality of Sleep'], hue = 'Gender', height = 5)
plt.title('Factors affecting quality of sleep', pad = 20, fontsize = 15)
plt.axis('tight')
plt.show()
```

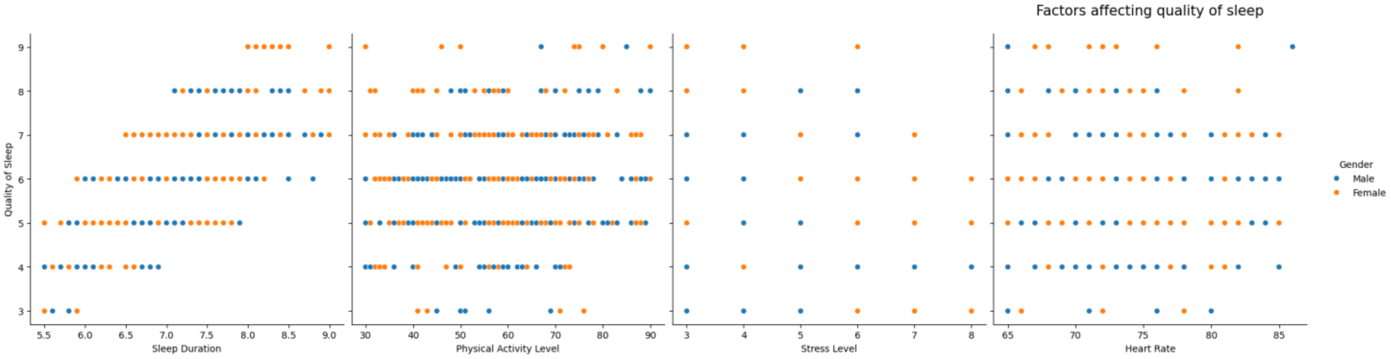

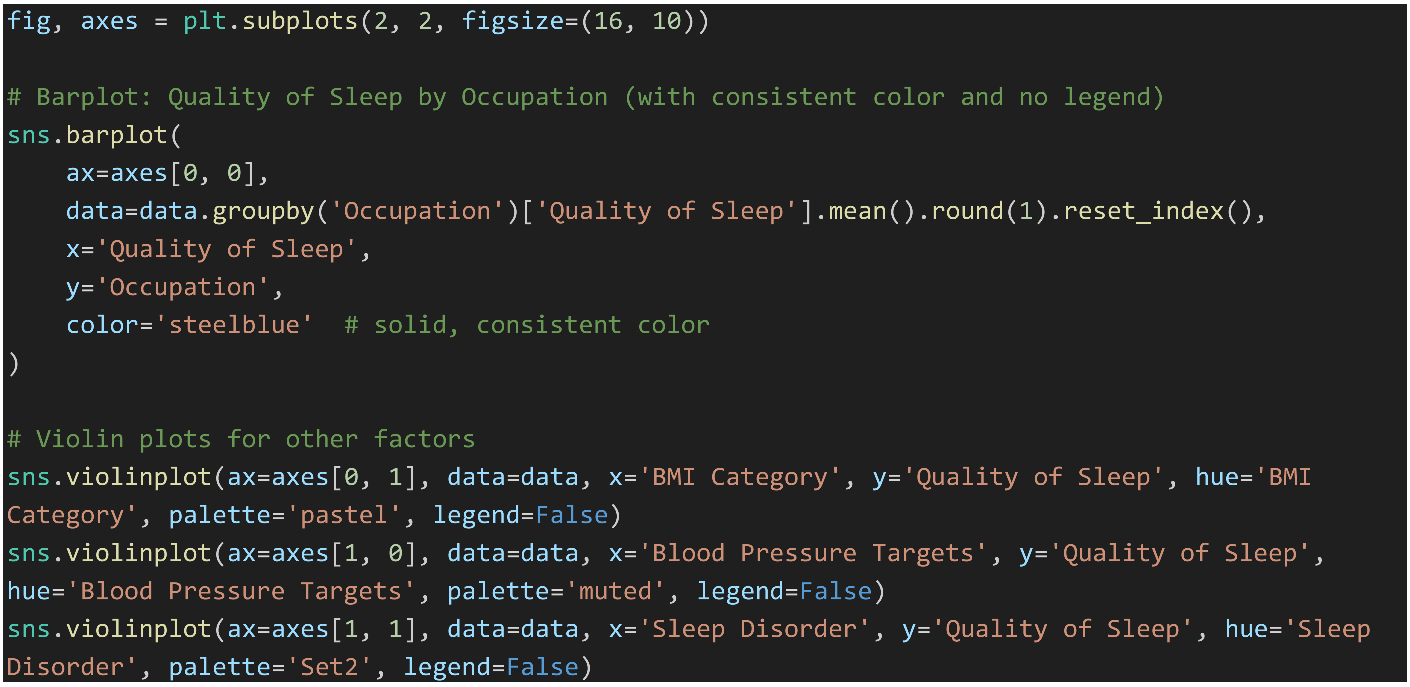

```
fig, axes = plt.subplots(2, 2, figsize=(16, 10))

# Barplot: Quality of Sleep by Occupation (with consistent color and no legend)
sns.barplot(
    ax=axes[0, 0],
    data=data.groupby('Occupation')['Quality of Sleep'].mean().round(1).reset_index(),
    x='Quality of Sleep',
    y='Occupation',
    color='steelblue'  # solid, consistent color
)

# Violin plots for other factors
sns.violinplot(ax=axes[0, 1], data=data, x='BMI Category', y='Quality of Sleep', hue='BMI
Category', palette='pastel', legend=False)
sns.violinplot(ax=axes[1, 0], data=data, x='Blood Pressure Targets', y='Quality of Sleep',
hue='Blood Pressure Targets', palette='muted', legend=False)
sns.violinplot(ax=axes[1, 1], data=data, x='Sleep Disorder', y='Quality of Sleep', hue='Sleep
Disorder', palette='Set2', legend=False)
```

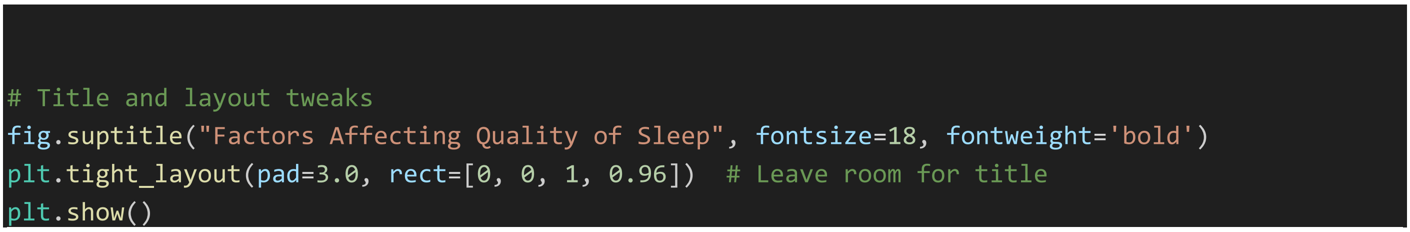

```
# Title and layout tweaks
fig.suptitle("Factors Affecting Quality of Sleep", fontsize=18, fontweight='bold')
plt.tight_layout(pad=3.0, rect=[0, 0, 1, 0.96])  # Leave room for title
plt.show()
```

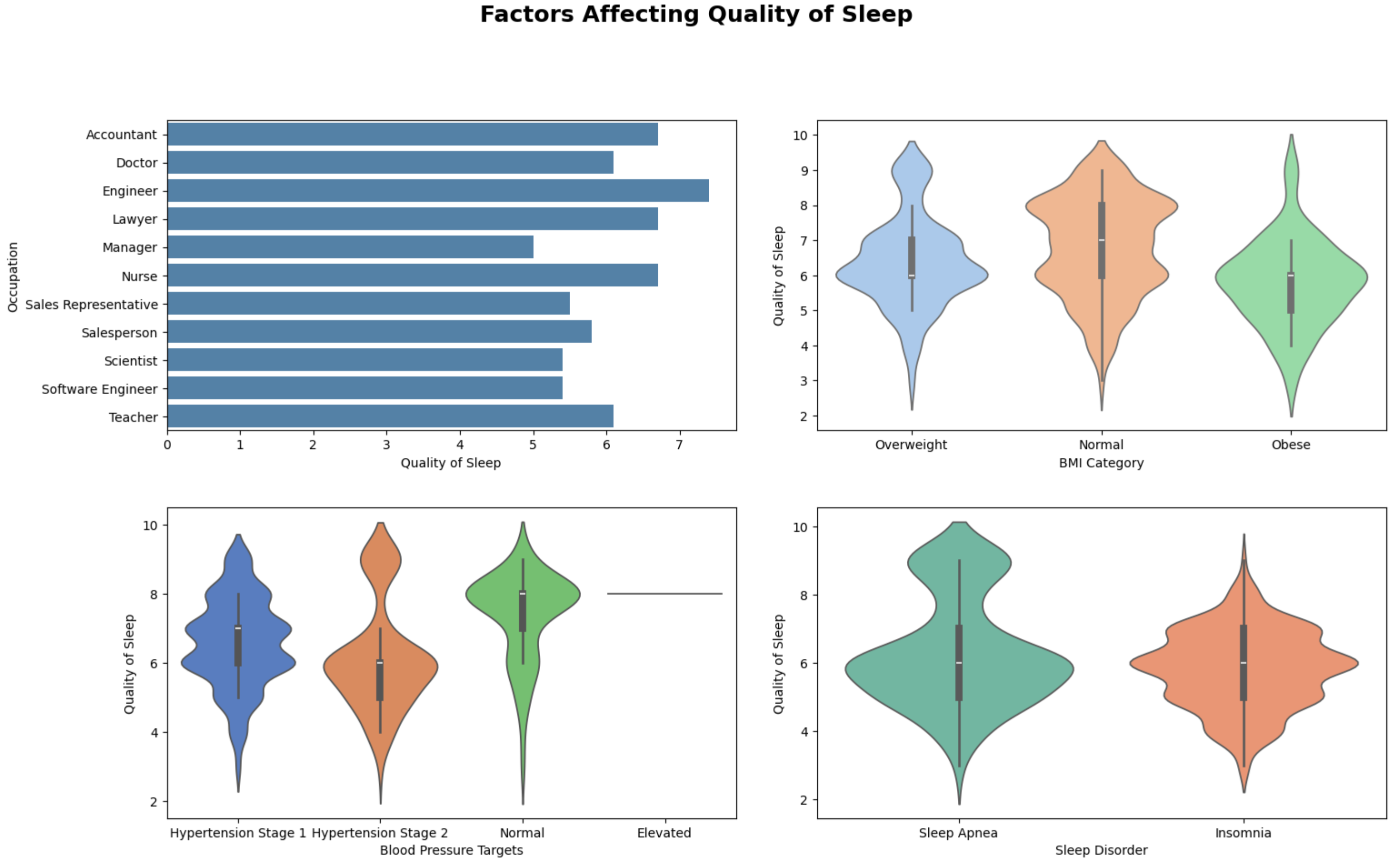

2. Data Preprocessing

Data preprocessing was implemented as a comprehensive pipeline to ensure consistency and reproducibility:

a. Handling Missing Data

- Created categorical age groups from continuous age values (Child, Young Adults, Middle-aged Adults, Old Adults)
- Generated derived clinical categories including Blood Pressure Targets (Normal, Elevated, Hypertension Stage 1, etc.)
- Developed Heart Rate Targets classification (Bradycardia, Normal, Tachycardia)
- Standardized BMI category nomenclature ('Normal Weight' → 'Normal')

b. Feature Transformation

- Numerical features, such as age and serum cholesterol, were examined for outliers using box plots and the Interquartile Range (IQR) method. Outliers beyond 1.5 times the IQR from the lower and upper quartiles were replaced with the respective column median.

c. Feature Scaling and Normalization

- Applied min-max scaling to Physical Activity Level, Heart Rate, and Daily Steps
- Implemented Box-Cox transformation for Sleep Duration, Quality of Sleep, and Stress Level to achieve better distributional properties

d. Advanced Feature Creation

- means clustering (10 clusters) on age, stress level, heart rate, physical activity, and daily steps
- Applied Principal Component Analysis (PCA) to capture complex interactions between numerical features

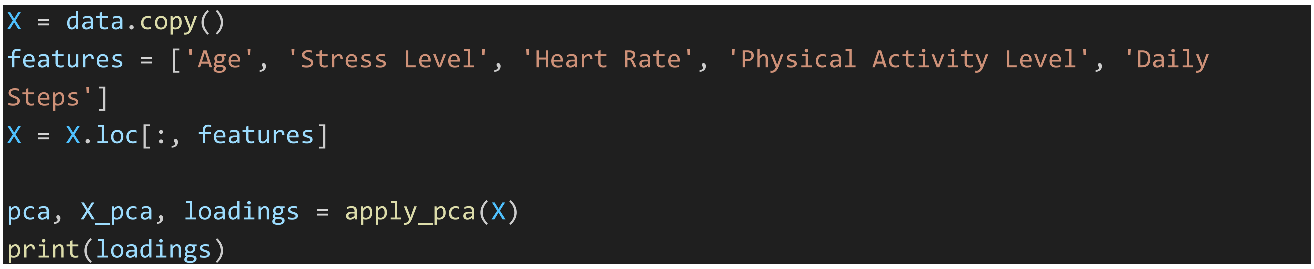

```
X = data.copy()
features = ['Age', 'Stress Level', 'Heart Rate', 'Physical Activity Level', 'Daily
Steps']
X = X.loc[:, features]

pca, X_pca, loadings = apply_pca(X)
print(loadings)
```

- Evaluated feature importance using mutual information scores to identify the most predictive variables

3. Model Development and Training

We employed a systematic approach to develop regression models for predicting sleep quality:

a. Model Selection

- Decision Tree Regressor as the primary model due to its interpretability
- Random Forest Regressor for comparison and potential performance improvement

b. Training Methodology

- Implemented an 80-20 train-test split with random stratification

- Created comprehensive preprocessing pipelines (SimpleImputer for numerical features, OneHotEncoder for categorical variables)

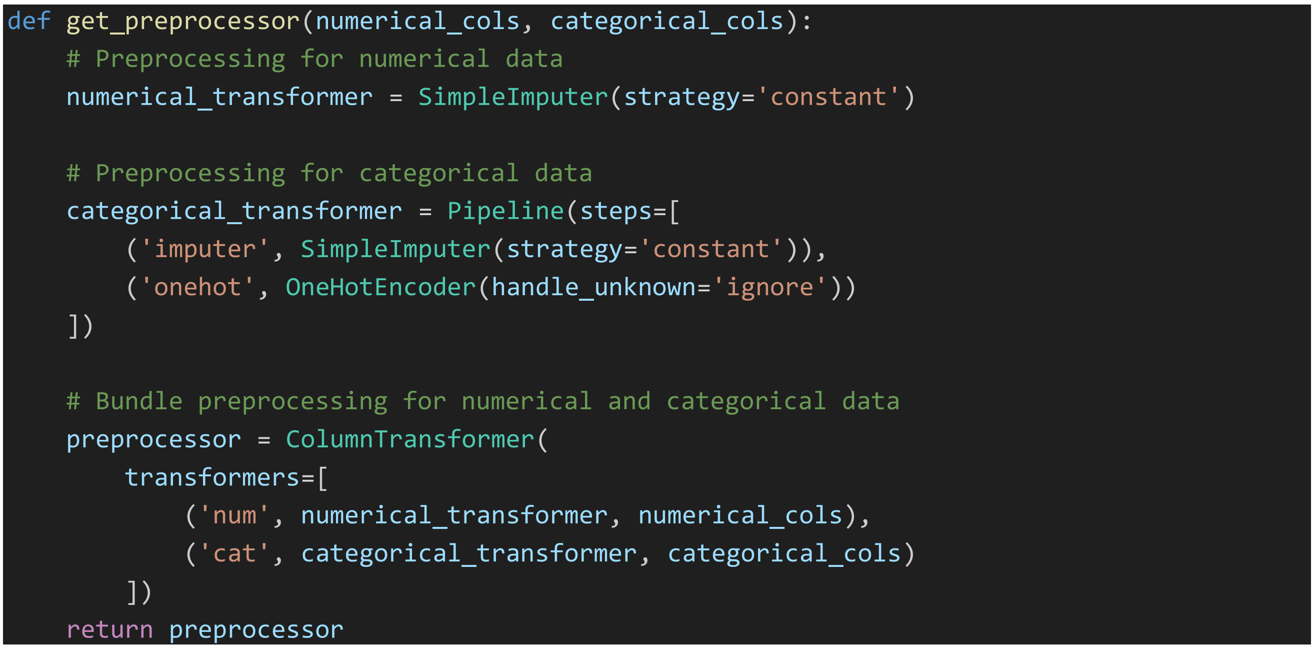

```python
def get_preprocessor(numerical_cols, categorical_cols):
    # Preprocessing for numerical data
    numerical_transformer = SimpleImputer(strategy='constant')

    # Preprocessing for categorical data
    categorical_transformer = Pipeline(steps=[
        ('imputer', SimpleImputer(strategy='constant')),
        ('onehot', OneHotEncoder(handle_unknown='ignore'))
    ])

    # Bundle preprocessing for numerical and categorical data
    preprocessor = ColumnTransformer(
        transformers=[
            ('num', numerical_transformer, numerical_cols),
            ('cat', categorical_transformer, categorical_cols)
        ])
    return preprocessor
```

- Employed 5-fold cross-validation to ensure model stability and generalizability

4. Hyperparameter Tuning and Optimization

To maximize prediction accuracy, we conducted systematic hyperparameter optimization:

*a. **Decision Tree Optimization***

- *Evaluated multiple complexity parameters with max_leaf_nodes: [5, 25, 50, 100, 250, 500, 1000, 5000]*

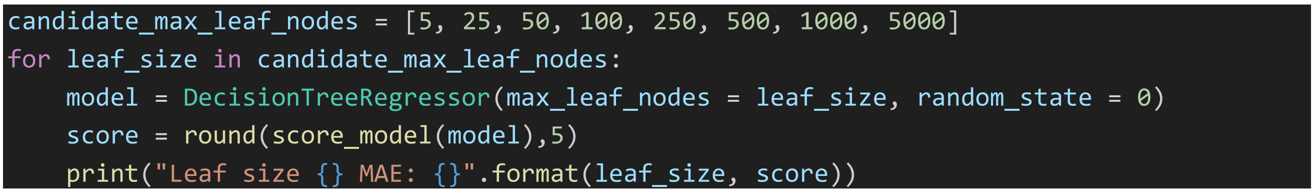

```python
candidate_max_leaf_nodes = [5, 25, 50, 100, 250, 500, 1000, 5000]
for leaf_size in candidate_max_leaf_nodes:
    model = DecisionTreeRegressor(max_leaf_nodes = leaf_size, random_state = 0)
    score = round(score_model(model),5)
    print("Leaf size {} MAE: {}".format(leaf_size, score))
```

- *Selected optimal complexity based on Mean Absolute Error minimization*

*b. **Random Forest Tuning***

- *Tested various ensemble configurations:*
    - *Different n_estimators values (50, 100, 200)*
    - *Alternative split criteria (absolute_error)*

   - *Various tree constraints (min_samples_split=20, max_depth=7)*

c. ***Performance Assessment***

- *Calculated Mean Absolute Error (MAE) for each model configuration*
- *Identified Decision Tree with max_leaf_nodes=25 as the optimal model (MAE: 0.24157)*

5. Advanced Prediction System Development

We extended beyond traditional machine learning by developing a hybrid prediction system:

a. Domain Knowledge Integration

- Created a comprehensive sleep prediction algorithm incorporating 10+ domain-specific factors
- Implemented age-specific, activity-level, and stress-based adjustments
- Added personalized factors for BMI, sleep disorders, heart rate, daily steps, gender, and occupation

b. **Personalized Recommendation Engine**

- Developed condition-specific sleep recommendations based on individual health profiles
- Created algorithms to prioritize recommendations based on user-specific risk factors
- Implemented context-aware suggestion filtering to ensure most relevant guidance

6. Evaluation and Visualization

The final system was evaluated using multiple approaches:

a. **Performance Metrics**

- Mean Absolute Error for regression accuracy assessment
- Comparative analysis between pure ML approach and hybrid system

b. **Interactive Visualization Suite**

- Developed multi-perspective interactive dashboard using streamlit
- Created personalized PDF report generation system
- Implemented health metrics assessment with radar charts

This methodology enabled us to develop a comprehensive sleep quality prediction and recommendation system that balances technical accuracy with practical, personalized guidance for improving sleep health.

Results

Model Performance

Model	MAE	Hyperparameter Optimization
Decision Tree (max_leaf_nodes=25)	0.42	Best performing leaf size among 8 tested configurations
Random Forest (n_estimators = 200, min_samples_split = 20)	0.43	Best performing among 5 tested configurations

Conclusion

1. The Decision Tree Regressor (max_leaf_nodes=25) achieved optimal performance with MAE of 0.42, while Random Forest models provided consistent results (MAE ~0.43), demonstrating the effectiveness of tree-based methods for sleep quality prediction.
2. Key predictors—Physical Activity Level, Stress Level, Sleep Duration, Heart Rate, and Daily Steps—align with established sleep medicine research, validating our feature engineering approach and mutual information analysis findings.
3. The hybrid prediction system combining machine learning with domain knowledge adjustments for 10 individual factors (including age, activity level, stress, and physiological metrics) enables personalized sleep recommendations that address specific health profiles and lifestyle patterns.
4. Future work: integrate continuous sleep monitoring data, expand the model to account for seasonal and environmental factors, incorporate cognitive performance metrics, and develop smartphone integration for real-time sleep optimization recommendations.

References

[1] Hirshkowitz, M., et al., "National Sleep Foundation's sleep time duration recommendations: methodology and results summary," Sleep Health, vol. 1, no. 1, pp. 40-43, 2015.

[2] Buysse, D.J., "Sleep health: can we define it? Does it matter?," Sleep, vol. 37, no. 1, pp. 9-17, 2014.

[3] Watson, N.F., et al., "Recommended amount of sleep for a healthy adult: a joint consensus statement of the American Academy of Sleep Medicine and Sleep Research Society," Sleep, vol. 38, no. 6, pp. 843-844, 2015.

[4] Breiman, L., "Random forests," Machine Learning, vol. 45, no. 1, pp. 5-32, 2001.

[5] Breiman, L., et al., "Classification and Regression Trees," Wadsworth, Belmont, CA, 1984.

[6] Chen, T. and Guestrin, C., "XGBoost: A scalable tree boosting system," in Proc. 22nd ACM SIGKDD, pp. 785-794, 2016.

[7] Pedregosa, F., et al., "Scikit-learn: Machine learning in Python," Journal of Machine Learning Research, vol. 12, pp. 2825-2830, 2011.

[8] Box, G.E.P. and Cox, D.R., "An analysis of transformations," Journal of the Royal Statistical Society B, vol. 26, no. 2, pp. 211-252, 1964.

[9] Kraskov, A., et al., "Estimating mutual information," Physical Review E, vol. 69, no. 6, 066138, 2004.[10] Cappuccio, F.P., et al., "Sleep duration and all-cause mortality: a systematic review and meta-analysis of prospective studies," Sleep, vol. 33, no. 5, pp. 585-592, 2010.